

# Data Structures and Algorithms:

## Exercise 1: E-commerce Search – Linear vs Binary Search

### 1. Theory

#### 1. Problem Context

Quickly locating an item ("Product") in a large catalog—critical for user experience and performance.

#### 2. Linear Search

- **Description:** Scan one element at a time until you find the target (or reach the end).
- **Use When:** Data is unsorted, or catalog is small.
- **Complexity**
  - Best Case:  $O(1)$  (item at index 0)
  - Average/Worst Case:  $O(n)$

#### 3. Binary Search

- **Description:** Repeatedly split a *sorted* array in half, compare mid-point to target, then recurse or iterate on the appropriate half.
- **Prerequisite:** Data must be sorted by key (here, `productId`).
- **Complexity**
  - Best/Average/Worst:  $O(\log n)$

#### 4. Asymptotic Notation & Trade-offs

- Sorting an array:  $O(n \log n)$  up-front cost for binary search, but each lookup is  $O(\log n)$ .
- Linear search: no preprocessing, but each lookup costs  $O(n)$ .

## 2. Code Summary

- We define a simple `Product` class (`id`, `name`, `category`).
- `linearSearch(Product[] arr, int targetId)` loops from start to end.
- `binarySearch(Product[] sortedArr, int targetId)` implements the classic iterative algorithm.
- In `main`, we generate a large sample array, pick a random `targetId`, then time each search using `System.nanoTime()`.

## 3. Test Cases

Case	Data Setup	Expected Behavior
1. Found at Beginning	<code>products = [{id:10,...}, {id:20,...}, ...], targetId=10</code>	Linear: index 0 in few ns; Binary: sorted index small (exact middle may shift), but target exists.
2. Not Found	<code>products = [{1,...},{2,...},{3,...}], targetId=99</code>	Both return <code>-1</code> —linear scans all; binary halves until empty.
3. Large Catalog	<code>n=1 000 000</code> , random <code>targetId</code>	Linear $\sim O(n)$ tens of millions of ns; binary $\sim O(\log n)$ (<50 comparisons).
4. Worst-Case for Binary	Already sorted, <code>targetId</code> outside range	Binary returns <code>-1</code> in $\sim \log n$ steps.

## 4. Code Output:

```
Searching for Product ID: 32134
Linear -> index=32133, time=337,200 ns
Binary -> index=32133, time=18,000 ns

Big-O Summary:
  Linear Search:  O(n)
  Binary Search:  O(log n)  (requires sorted data)

Process finished with exit code 0
```

# Exercise 2: Financial Forecasting – Recursive vs Iterative vs Memoized

## 1. Theory

### 1. Recursion

- **Definition:** A method calls itself on a smaller portion until a base case.
- **Advantages:** Declarative, mirrors mathematical definitions (e.g.,  $FV_t = FV_{t-1} \cdot (1+r_t)$ ).
- **Drawbacks:**  $O(n)$  stack space, potential for stack overflow.

### 2. Iterative Approach

- Loops over rates, updating a running total.
- $O(n)$  time,  $O(1)$  extra space.

### 3. Memoization

- Cache results of sub-calls to avoid recomputation.
- Useful when you need multiple overlapping queries (e.g., “What’s year 2?” and later “What’s year 4?”).

### 4. Complexity

- **Pure Recursion:** Time  $O(n)$ , Space  $O(n)$  (stack depth).
- **Iterative:** Time  $O(n)$ , Space  $O(1)$ .
- **Memoized:** Time  $O(n)$  for first full run, then  $O(1)$  per cache hit; Space  $O(n)$  for cache + stack.

## 2. Code Summary

- `forecastRecursive(base, rates, idx)` calls itself down to index  $-1$ , then unwinds multiplying by  $(1 + \text{rates}[i])$ .
- `forecastIterative(base, rates)` loops once.
- `forecastMemo(base, rates, idx, cache)` checks `cache` before recursing, stores results.
- In `main`, we
  1. Run each method and print the forecasted value.
  2. Track and print `maxDepth` for recursion.
  3. Simulate **repeated queries** for Year 2 and Year 4 to show memo's speedup.
  4. Print a complexity analysis.

## 3. Test Cases

Case	Input	Expected Notes
1. Uniform Rates	<code>base=1000</code> , <code>rates=[0.05,0.05,0.05]</code>	All methods $\rightarrow 1000 \cdot 1.05^3 \approx 1157.63$ . Recursion depth = 4; iterative uses no extra stack; memo builds cache of size 3.
2. Zero Growth	<code>base=500</code> , <code>rates=[0,0,0]</code>	All methods $\rightarrow 500$ . Tests base case behavior and multiplication-by-1.
3. Repeated Queries	Query Year 1 then Year 3	Iterative: full loop twice (2 ops + 3 ops). Memo: first Year 1 populates <code>cache[1]</code> , then Year 3 reuses Year 1 result.
4. Single Year	<code>rates=[0.1]</code>	Ensures base case ( <code>idx &lt; 0</code> ) and single recursion step work ("bankruptcy" edge-case style).

## 4. Code Output

```
=== Financial Forecast Demo ===
Base amount: ₹1000.0
Rates (annual): [0.07, 0.06, 0.065, 0.055]

1) Recursive result: ₹1274.36 (max stack depth = 5)
2) Iterative result: ₹1274.36 (no extra stack)
3) Memoized recursion: ₹1274.36 (cache entries = 4)

--- Test: Repeated Queries (Years 2 & 4) ---
Iterative: Year 2 in 5,000 ns; Year 4 in 3,100 ns
MemoRec : Year 2 in 5,700 ns; Year 4 in 6,600 ns

→ You'll see MemoRec for Year 4 is faster than recomputing iteratively from scratch.

=== Complexity Analysis ===
Pure Recursion:
  • Time:  $O(n)$  calls
  • Space:  $O(n)$  stack frames (maxDepth =  $n+1$ )
Iteration:
  • Time:  $O(n)$  loop
  • Space:  $O(1)$ 
Memoized Recursion:
  • Time:  $O(n)$  overall, but re-queries in  $O(1)$ 
  • Space:  $O(n)$  for cache +  $O(n)$  stack (can be reduced with tail recursion elimination)
```

# Design Principles:

## Exercise 1: Singleton Pattern – Thread-Safe Logger

### 1. Theory

#### 1. Intent

Ensure only one instance of a class exists, provide a global access point.

#### 2. Common Pitfalls

- **Non-thread-safe** lazy instantiation can create multiple instances under concurrency.
- **Eager instantiation** (static field) solves safety but may waste resources.

#### 3. Double-Checked Locking

- Uses a `volatile` instance and checks `if (instance==null)` both outside and inside a `synchronized` block.
- Ensures thread-safe lazy initialization with minimal synchronization overhead.

#### 4. SOLID Principles

- **SRP**: `Logger` has one responsibility—recording messages.
- **DIP**: Code depends on the `Logger` abstraction, not on how it's instantiated.

### 2. Code Summary

- `Logger` is a **static nested class** inside `SingletonDemo`.

- `private static volatile Logger instance;`
- `private Logger()` prevents external instantiation.
- `public static Logger getInstance()` implements double-checked locking.
- `log(level,msg)` prints a timestamped message.
- In `main`, we:
  1. Obtain two references, demonstrate `==`.
  2. Spawn multiple threads that each call `getInstance().log(...)`.
  3. Show only one “Logger initialized” line, proving singleton.
  4. Print analysis.

### 3. Test Cases

Case	Action	Expected Outcome
1. Single-Thread Instantiation	Call <code>getInstance()</code> twice	Both refs are <code>==</code> ; constructor prints once.
2. Multi-Thread Race	Spawn 3 threads simultaneously	Only one thread should actually construct; others see non-null instance.
3. Reflection Attack (Discussion)	Explain why private ctor + no reflection-safe code can still be broken by reflection—mention <code>enum</code> singletons.	
4. Serialization (Discussion)	Note: to guard against multiple instances via serialization, implement <code>readResolve()</code> returning <code>instance</code> .	

## 4. Code Output

```
=== Singleton Logger Demo ===
Logger initialized at 2025-06-22 13:52:54
[2025-06-22 13:52:54] INFO First message
Same instance? true

Spawning threads:
[2025-06-22 13:52:54] DEBUG log from T1
[2025-06-22 13:52:54] DEBUG log from T2
[2025-06-22 13:52:54] DEBUG log from T3

Still same? true

--- Analysis ---
1) Private ctor → no outside instantiation.
2) Double-checked locking → thread-safe.
3) Only one "Logger initialized" printed → singleton holds.

Process finished with exit code 0
```



# Exercise 2: Factory Method Pattern – Document Management

## 1. Theory

### 1. Intent

Define an interface for creating an object, but let subclasses decide which class to instantiate.

### 2. Structure

- **Product:** `Document` interface.
- **ConcreteProducts:** `WordDocument`, `PdfDocument`, `ExcelDocument`.
- **Creator:** abstract `DocumentFactory` with `createDocument()`.
- **ConcreteCreators:** `WordDocumentFactory`, `PdfDocumentFactory`, `ExcelDocumentFactory`.

### 3. Advantages

- Decouples client code from concrete classes.
- Easy to add new document types without modifying existing client logic.
- Adheres to **OCP**: open for extension (new factories), closed for modification.

## 2. Code Summary

- Define `interface Document { open(); printInfo(); }`.
- Implement three document types that simulate real metadata (random page counts, word counts, sheet counts).
- Abstract `DocumentFactory` declares `createDocument(String filename)`.
- Concrete factories override to return their specific document.
- Client (`Main`) collects file names, chooses a factory via extension, calls `createDocument()`, then `open()` & `printInfo()`.

## 3. Test Cases

Case	Files List	Expected Behavior
1. Supported Types	<code>["Report.docx", "Data.xlsx", "Slides.pdf"]</code>	Each recognized, opens appropriately, prints metadata.
2. Unsupported Type	<code>["Notes.txt", "Image.png"]</code>	Prints "Skipping unsupported file type" and does not attempt <code>open()</code> .
3. Mixed Case Extensions	<code>["report.PdF", "data.XLsX"]</code>	Lower-casing in factory lookup ensures these still get handled.
4. Empty Filename or No Extension	<code>["README", ""]</code>	Factory returns <code>null</code> , client skips—tests robustness.

## 4. Code Output

```
=== Factory Method Demo ===

Opening Word document: Report.docx
WordDoc: name="Report.docx", words=858

Opening PDF document: Slides.pdf
PdfDoc: name="Slides.pdf", pages=43

Opening Excel document: Data.xlsx
ExcelDoc: name="Data.xlsx", sheets=3

Skipping unsupported file type: Notes.txt

Opening Excel document: Budget.xlsx
ExcelDoc: name="Budget.xlsx", sheets=1

Opening PDF document: Summary.pdf
PdfDoc: name="Summary.pdf", pages=26
```